hw3

Jiajian Huang 4/6/2019

Q8 CHP 9

```
library(caret)
##
## Attaching package: 'caret'
## The following object is masked from 'package:purrr':
##
##
       lift
## The following object is masked from 'package:mosaic':
##
##
       dotPlot
## The following object is masked from 'package:survival':
##
##
       cluster
set.seed(2019)
train = sample(dim(OJ)[1], 800)
OJ.train = OJ[train, ]
OJ.test = OJ[-train, ]
```

```
library(e1071)
```

```
##
## Attaching package: 'e1071'
```

```
## The following object is masked from 'package:Hmisc':
##
## impute
```

```
svm.linear <- svm(Purchase ~ ., data = OJ.train, kernel = "linear", cost = 0.01)
summary(svm.linear)</pre>
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = OJ.train, kernel = "linear",
##
       cost = 0.01)
##
##
## Parameters:
##
                C-classification
      SVM-Type:
    SVM-Kernel:
##
                 linear
##
                 0.01
          cost:
                 0.0555556
##
         gamma:
##
## Number of Support Vectors:
##
    (222219)
##
##
##
## Number of Classes:
##
## Levels:
##
    CH MM
```

From the output of R's summary function we can see that 441 observations are used as support vector. Moreover, the support vectors are almost equally split among the classes.

C.

```
train.pred <- predict(svm.linear, OJ.train)
table(OJ.train$Purchase, train.pred)</pre>
```

```
## train.pred
## CH MM
## CH 447 55
## MM 80 218
```

```
(80 + 55)/ ( 447+80+55+218)
```

```
## [1] 0.16875
```

The training error rate is 0.1688.

```
test.pred <- predict(svm.linear, OJ.test)
table(OJ.test$Purchase, test.pred)</pre>
```

```
## test.pred
## CH MM
## CH 135 16
## MM 26 93
```

```
(26+16)/(135+16+26+93)
```

```
## [1] 0.155556
```

The testing error rate is 0.1556.

d

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##
    cost
##
      10
##
## - best performance: 0.175
##
## - Detailed performance results:
##
             cost
                    error dispersion
## 1
       0.01000000 0.17875 0.04084609
##
       0.01778279 0.17875 0.04126894
       0.03162278 0.17875 0.04291869
##
   3
##
       0.05623413 0.18250 0.04133199
##
   5
       0.10000000 0.18125 0.04419417
       0.17782794 0.18125 0.04419417
##
## 7
       0.31622777 0.18000 0.04495368
##
       0.56234133 0.17875 0.04041881
  8
## 9
       1.00000000 0.17875 0.04041881
       1.77827941 0.18000 0.03827895
## 10
## 11
       3.16227766 0.18000 0.03827895
## 12
       5.62341325 0.17625 0.03606033
## 13 10.00000000 0.17500 0.03435921
```

Tuning shows that optimal cost is 3.162278.

e.

```
svm.linear <- svm(Purchase ~ ., kernel = "linear", data = OJ.train, cost = tune.out$b
est.parameter$cost)
train.pred <- predict(svm.linear, OJ.train)
table(OJ.train$Purchase, train.pred)</pre>
```

```
## train.pred
## CH MM
## CH 441 61
## MM 75 223
```

```
(75+61)/(441 + 61+ 75+223)
```

```
## [1] 0.17
```

The training error rate is now 17%

```
test.pred <- predict(svm.linear, OJ.test)
table(OJ.test$Purchase, test.pred)</pre>
```

```
## test.pred
## CH MM
## CH 135 16
## MM 25 94
```

```
(25+16)/(135+25+16+94)
```

```
## [1] 0.1518519
```

The test error rate is 15.18%.

f.

```
svm.radial <- svm(Purchase ~ ., kernel = "radial", data = OJ.train)
summary(svm.radial)</pre>
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = OJ.train, kernel = "radial")
##
##
## Parameters:
##
      SVM-Type: C-classification
##
    SVM-Kernel:
                radial
##
          cost:
                 1
##
         gamma:
                 0.0555556
##
## Number of Support Vectors:
                                368
##
    (185 183)
##
##
##
## Number of Classes: 2
##
## Levels:
##
    CH MM
```

```
train.pred <- predict(svm.radial, OJ.train)
table(OJ.train$Purchase, train.pred)</pre>
```

```
## train.pred
## CH MM
## CH 465 37
## MM 80 218
```

```
(31+18)/(133+18+31+88)
```

```
## [1] 0.1814815
```

```
test.pred <- predict(svm.radial, OJ.test)
table(OJ.test$Purchase, test.pred)</pre>
```

```
## test.pred
## CH MM
## CH 133 18
## MM 31 88
```

```
(80+37)/(465+37+80+218)
```

```
## [1] 0.14625
```

The classifier has a training error of 14.63% and a test error of 18.15%.

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
    cost
##
##
       1
##
## - best performance: 0.17125
##
## - Detailed performance results:
##
                    error dispersion
             cost
       0.01000000 0.37250 0.06368324
## 1
       0.01778279 0.37250 0.06368324
## 2
## 3
       0.03162278 0.36375 0.07008180
## 4
       0.05623413 0.21750 0.04338138
## 5
       0.10000000 0.17875 0.05560588
## 6
       0.17782794 0.17125 0.04715886
       0.31622777 0.17375 0.04059026
## 7
## 8
       0.56234133 0.17375 0.04143687
       1.00000000 0.17125 0.04450733
## 9
       1.77827941 0.17250 0.04401704
## 10
       3.16227766 0.17750 0.04158325
## 11
       5.62341325 0.17375 0.04185375
## 12
## 13 10.00000000 0.18125 0.04497299
```

Tuning shows that optimal cost is 0.17.

g.

```
svm.poly <- svm(Purchase ~ ., kernel = "polynomial", data = OJ.train, degree = 2)
summary(svm.poly)</pre>
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = OJ.train, kernel = "polynomial",
       degree = 2)
##
##
##
## Parameters:
##
      SVM-Type:
                 C-classification
##
                 polynomial
    SVM-Kernel:
##
                  1
          cost:
##
        degree:
                  2
##
                  0.0555556
         gamma:
##
        coef.0:
                  0
##
## Number of Support Vectors:
                                 435
##
##
    (222 213)
##
##
## Number of Classes: 2
##
## Levels:
##
    CH MM
train.pred <- predict(svm.poly, OJ.train)</pre>
table(OJ.train$Purchase, train.pred)
```

```
## train.pred
## CH MM
## CH 470 32
## MM 109 189
```

```
(109+32)/(470+32+109+189)
```

```
## [1] 0.17625
```

```
test.pred <- predict(svm.poly, OJ.test)
table(OJ.test$Purchase, test.pred)</pre>
```

```
## test.pred

## CH MM

## CH 134 17

## MM 41 78
```

```
(41+17)/(134+17+41+78)
```

```
## [1] 0.2148148
```

With polynomial kernel degree=2, the classifier has a training error of 17.63% and a test error of 21.48%.

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##
    cost
##
      10
##
## - best performance: 0.18125
##
## - Detailed performance results:
##
                    error dispersion
             cost
## 1
       0.01000000 0.37125 0.06402311
## 2
       0.01778279 0.35125 0.05905800
       0.03162278 0.34250 0.06671873
##
   3
       0.05623413 0.31000 0.06609127
##
   5
       0.10000000 0.29000 0.06687468
##
##
       0.17782794 0.25000 0.06614378
##
   7
       0.31622777 0.20875 0.03821086
       0.56234133 0.19625 0.04332131
## 8
       1.00000000 0.19500 0.03343734
##
   9
       1.77827941 0.18500 0.04594683
## 10
## 11
       3.16227766 0.18375 0.03488573
## 12
       5.62341325 0.18625 0.04016027
## 13 10.00000000 0.18125 0.04379958
```

Tuning shows that optimal cost is 0.17.

```
svm.poly <- svm(Purchase ~ ., kernel = "polynomial", degree = 2, data = OJ.train, cos
t = tune.out$best.parameter$cost)
summary(svm.poly)</pre>
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = OJ.train, kernel = "polynomial",
##
       degree = 2, cost = tune.out$best.parameter$cost)
##
##
## Parameters:
##
      SVM-Type:
                 C-classification
##
    SVM-Kernel:
                 polynomial
##
                 10
          cost:
##
                 2
        degree:
##
         gamma:
                 0.0555556
##
        coef.0:
##
## Number of Support Vectors:
##
##
    (173 163)
##
##
## Number of Classes: 2
##
## Levels:
    CH MM
##
train.pred <- predict(svm.poly, OJ.train)</pre>
table(OJ.train$Purchase, train.pred)
```

```
## train.pred

## CH MM

## CH 466 36

## MM 79 219
```

```
(79+36)/(466+36+79+219)
```

```
## [1] 0.14375
```

```
test.pred <- predict(svm.poly, OJ.test)
table(OJ.test$Purchase, test.pred)</pre>
```

```
## test.pred

## CH MM

## CH 131 20

## MM 33 86
```

```
(33+20)/(131+20+33+86)
```

```
## [1] 0.1962963
```

The classifier has a training error of 14.38% and a test error of 19.63%.

h.

##

grow

The default gamma approach provides the best results on this data.

Dataset Khan

random forests

```
data("Khan")
library(randomForest)
## randomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:dplyr':
##
##
       combine
  The following object is masked from 'package:ggplot2':
##
##
       margin
   The following object is masked from 'package:gridExtra':
##
##
##
       combine
   The following object is masked from 'package:imager':
##
```

```
## [1]
         63 2308
dim(Khan$xtest)
## [1]
         20 2308
table(Khan$ytrain)
##
## 1 2 3 4
## 8 23 12 20
table(Khan$ytest)
##
## 1 2 3 4
## 3 6 6 5
khan_train = data.frame(x = Khan$xtrain, y = as.factor(Khan$ytrain))
khan test = data.frame(x = Khan$xtest, y = as.factor(Khan$ytest))
random = train(y ~ ., data = khan train, method = "rf", trControl = trainControl(meth
od = "cv"))
confusionMatrix(khan_train$y, predict(random, khan_train))
```

dim(Khan\$xtrain)

```
## Confusion Matrix and Statistics
##
##
            Reference
## Prediction 1
                  2
            1
##
                  0
##
            2
              0 23 0 0
##
            3
              0
                  0 12 0
##
              0
                  0 0 20
##
## Overall Statistics
##
##
                 Accuracy: 1
##
                    95% CI: (0.9431, 1)
      No Information Rate: 0.3651
##
##
      P-Value [Acc > NIR] : < 2.2e-16
##
##
                    Kappa: 1
##
   Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##
                       Class: 1 Class: 2 Class: 3 Class: 4
## Sensitivity
                           1.000
                                  1.0000
                                           1.0000
                                                     1.0000
## Specificity
                           1.000
                                  1.0000
                                           1.0000
                                                     1.0000
## Pos Pred Value
                                           1.0000
                           1.000
                                  1.0000
                                                     1.0000
## Neg Pred Value
                           1.000
                                  1.0000 1.0000
                                                     1.0000
## Prevalence
                           0.127
                                  0.3651
                                           0.1905
                                                     0.3175
## Detection Rate
                           0.127
                                  0.3651
                                           0.1905
                                                     0.3175
## Detection Prevalence
                           0.127
                                  0.3651
                                           0.1905
                                                     0.3175
## Balanced Accuracy
                           1.000
                                  1.0000
                                            1.0000
                                                     1.0000
```

```
confusionMatrix(khan_test$y, predict(random, khan_test))
```

```
## Confusion Matrix and Statistics
##
##
             Reference
## Prediction 1 2 3 4
##
            1 3 0 0 0
##
            2 0 6 0 0
            3 0 1 5 0
##
            4 0 0 0 5
##
##
## Overall Statistics
##
##
                   Accuracy: 0.95
                     95% CI: (0.7513, 0.9987)
##
##
       No Information Rate: 0.35
##
       P-Value [Acc > NIR] : 2.903e-08
##
##
                      Kappa : 0.932
##
    Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##
                         Class: 1 Class: 2 Class: 3 Class: 4
## Sensitivity
                             1.00
                                     0.8571
                                              1.0000
                                                          1.00
## Specificity
                                                          1.00
                             1.00
                                     1.0000
                                              0.9333
## Pos Pred Value
                             1.00
                                     1.0000
                                              0.8333
                                                          1.00
## Neg Pred Value
                             1.00
                                     0.9286
                                              1.0000
                                                          1.00
## Prevalence
                             0.15
                                     0.3500
                                              0.2500
                                                          0.25
## Detection Rate
                             0.15
                                     0.3000
                                              0.2500
                                                          0.25
## Detection Prevalence
                             0.15
                                     0.3000
                                              0.3000
                                                          0.25
## Balanced Accuracy
                             1.00
                                                          1.00
                                     0.9286
                                              0.9667
```

tree boosting

```
##
## Attaching package: 'xgboost'

## The following object is masked from 'package:dplyr':
##
## slice
```

```
library(gbm)
```

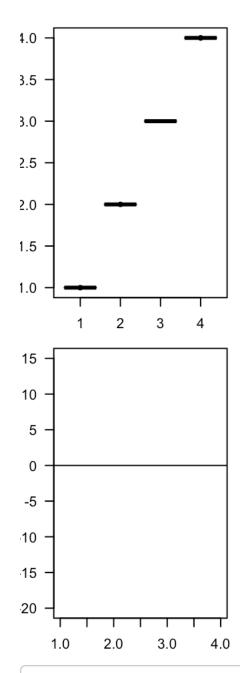
```
## Loaded gbm 2.1.5
```

```
set.seed(2019)
boost.khan = gbm(y ~., data = khan_train, distribution = "gaussian", n.trees = 5000)
boost.khan
```

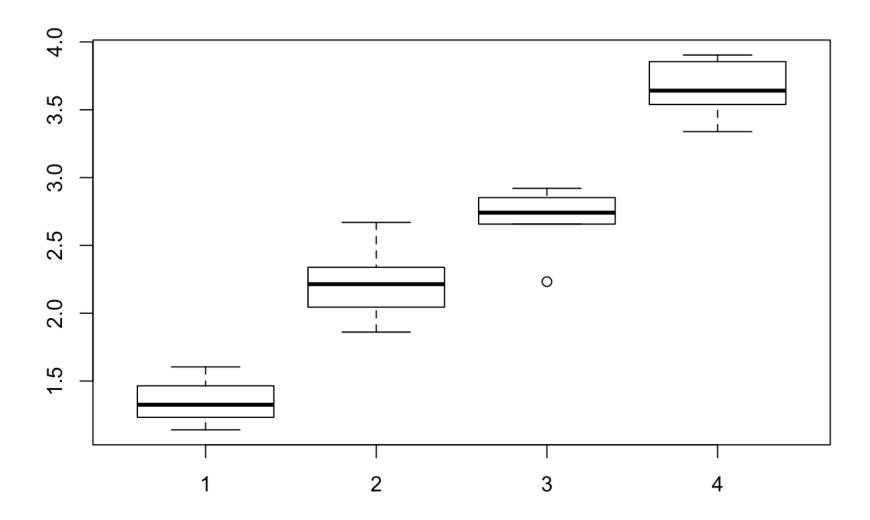
```
## gbm(formula = y ~ ., distribution = "gaussian", data = khan_train,
## n.trees = 5000)
## A gradient boosted model with gaussian loss function.
## 5000 iterations were performed.
## There were 2308 predictors of which 1845 had non-zero influence.
```

```
par(mfcol = c(2, 4), mar = c(2, 2, 1, 1), las = 1)
plot(khan_train$y, predict(boost.khan, n.trees = 5000))
plot(khan_train$y, predict(boost.khan, n.trees = 5000) - khan_train$y, ylim = c(-20, 15));abline(h = 0)
```

```
## Warning in Ops.factor(predict(boost.khan, n.trees = 5000), khan_train$y):
## '-' not meaningful for factors
```



```
plot(khan_test$y, predict(boost.khan, khan_test, n.tree = 500))
```



```
plot(khan_test$y, predict(boost.khan, khan_test, n.trees = 500) - khan_test$y, ylim = c(-20, 15));abline(h = 0)
```

```
## Warning in Ops.factor(predict(boost.khan, khan_test, n.trees = 500),
## khan_test$y): '-' not meaningful for factors
```

```
$\frac{9}{0} - \\
\frac{9}{0} - \\
\frac{9}{0} - \\
\frac{9}{0} - \\
\frac{9}{0} - \\
\frac{1}{0} - \\
\frac
```

```
set.seed(2019)
boost.han = gbm(y ~., data = khan_train, distribution = "gaussian", n.trees = 5000, i
nteraction.depth = 4)
```

```
yhat.btl.nn = predict(boost.khan, khan_test, n.trees = seq(500, 5000, 500))
dim(yhat.btl.nn)
```

```
## [1] 20 10
```

```
a = apply(yhat.bt1.nn, 2, function(x) mean(Khan$ytest - x)^2) round(a, 5)
```

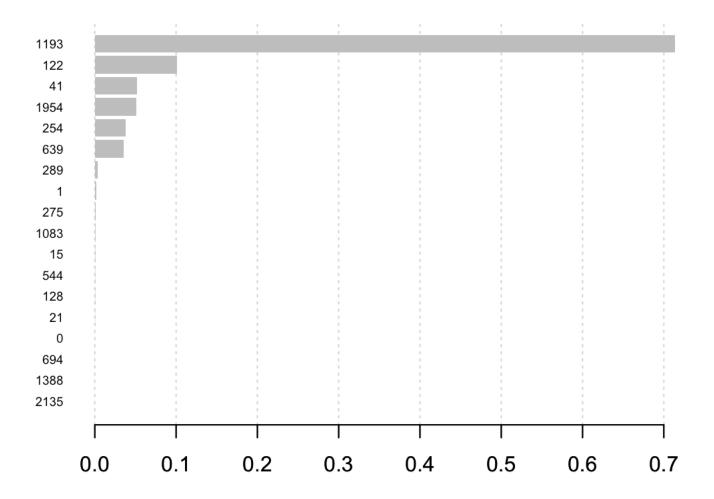
```
## 500 1000 1500 2000 2500 3000 3500 4000 4500
## 0.00340 0.00353 0.00353 0.00353 0.00353 0.00353 0.00353 0.00353
## 5000
## 0.00353
```

```
ound = 5, objective = "reg:linear")
## [1]
       train-rmse:0.386910
## [2] train-rmse:0.090095
## [3] train-rmse:0.032245
## [4] train-rmse:0.009076
## [5] train-rmse:0.003345
all.equal(as.numeric(predict(boost.x, Khan$xtest) > 0.5), Khan$ytest)
## [1] "Mean relative difference: 1.941176"
boost.y = xgboost(data = Khan$xtrain, label = Khan$ytrain, max depth = 1, eta = 0.1,
nround = 50, objective = "reg:linear")
## [1]
      train-rmse:2.213275
## [2] train-rmse:2.013805
## [3] train-rmse:1.835033
## [4] train-rmse:1.674052
## [5] train-rmse:1.529896
## [6] train-rmse:1.399689
## [7] train-rmse:1.281420
## [8] train-rmse:1.174630
## [9] train-rmse:1.078368
## [10] train-rmse:0.991666
## [11] train-rmse:0.912312
## [12] train-rmse:0.841279
## [13] train-rmse:0.776592
## [14] train-rmse:0.718592
## [15] train-rmse:0.665687
## [16] train-rmse:0.617682
## [17] train-rmse:0.574163
## [18] train-rmse:0.534266
## [19] train-rmse:0.498723
## [20] train-rmse:0.466249
## [21] train-rmse:0.437040
## [22] train-rmse:0.410511
## [23] train-rmse:0.386137
## [24] train-rmse:0.363850
## [25] train-rmse:0.343359
## [26] train-rmse:0.324529
## [27] train-rmse:0.307804
## [28] train-rmse:0.292361
## [29] train-rmse:0.278201
## [30] train-rmse:0.265033
```

boost.x = xgboost(data = Khan\$xtrain, label = Khan\$ytrain, max depth = 4, eta = 1, nr

```
## [31] train-rmse:0.253135
## [32] train-rmse:0.242135
## [33] train-rmse:0.231865
## [34] train-rmse:0.222444
## [35] train-rmse:0.213813
## [36] train-rmse:0.205722
## [37] train-rmse:0.198237
## [38] train-rmse:0.191119
## [39] train-rmse:0.184456
## [40] train-rmse:0.178209
## [41] train-rmse:0.172499
## [42] train-rmse:0.167112
## [43] train-rmse:0.161845
## [44] train-rmse:0.157029
## [45] train-rmse:0.152446
## [46] train-rmse:0.148043
## [47] train-rmse:0.143891
## [48] train-rmse:0.140015
## [49] train-rmse:0.136257
## [50] train-rmse:0.132729
```

```
importance_matrix = xgb.importance(model = boost.x)
xgb.plot.importance(importance_matrix = importance_matrix)
```



```
xgb.plot.importance(importance_matrix = xgb.importance(model =boost.y))
```

